

MAD Style: Multivalent Authorship Detection (MAD) Topic Models

David Dohan, Charles Marsh, Shubhro Saha, Max Simchowitz
Princeton University, Department of Computer Science

Goals

- Classify author writing style in a wide range of media.
- Extract compact representation of stylistic tendency.
- Determine which features are most indicative of writing style.

Introduction

In the *authorship detection* problem, one is given:

- A set of documents labeled (by author) on which to train.
- A set of anonymized documents to classify.

Methods for authorship detection have traditionally depended on lexical analysis of the text, making them relatively context-dependent.

Instead, the *MAD Topic Model* depends solely on syntactic and stylometric features (e.g., part-of-speech tags, meter), which are less context-dependent. MAD treats these features as vocabularies over which topic models can be determined, performing a Supervised Latent Dirichlet Allocation (SLDA) algorithm over n -gram stylistic features to determine authorship of anonymized text.

Preliminary results show significant improvement over more naive techniques (such as Logistic MLE) using the same features. As a by-product, MAD's topic models over the n -gram stylistic features can be used to extract compact representations of stylistic tendency and discern which features are most indicative of writing style.

Data

Our data.

Features

Our features.

Visualization

Replace this with some visualization.

Conclusion

Our (short) conclusion.

References

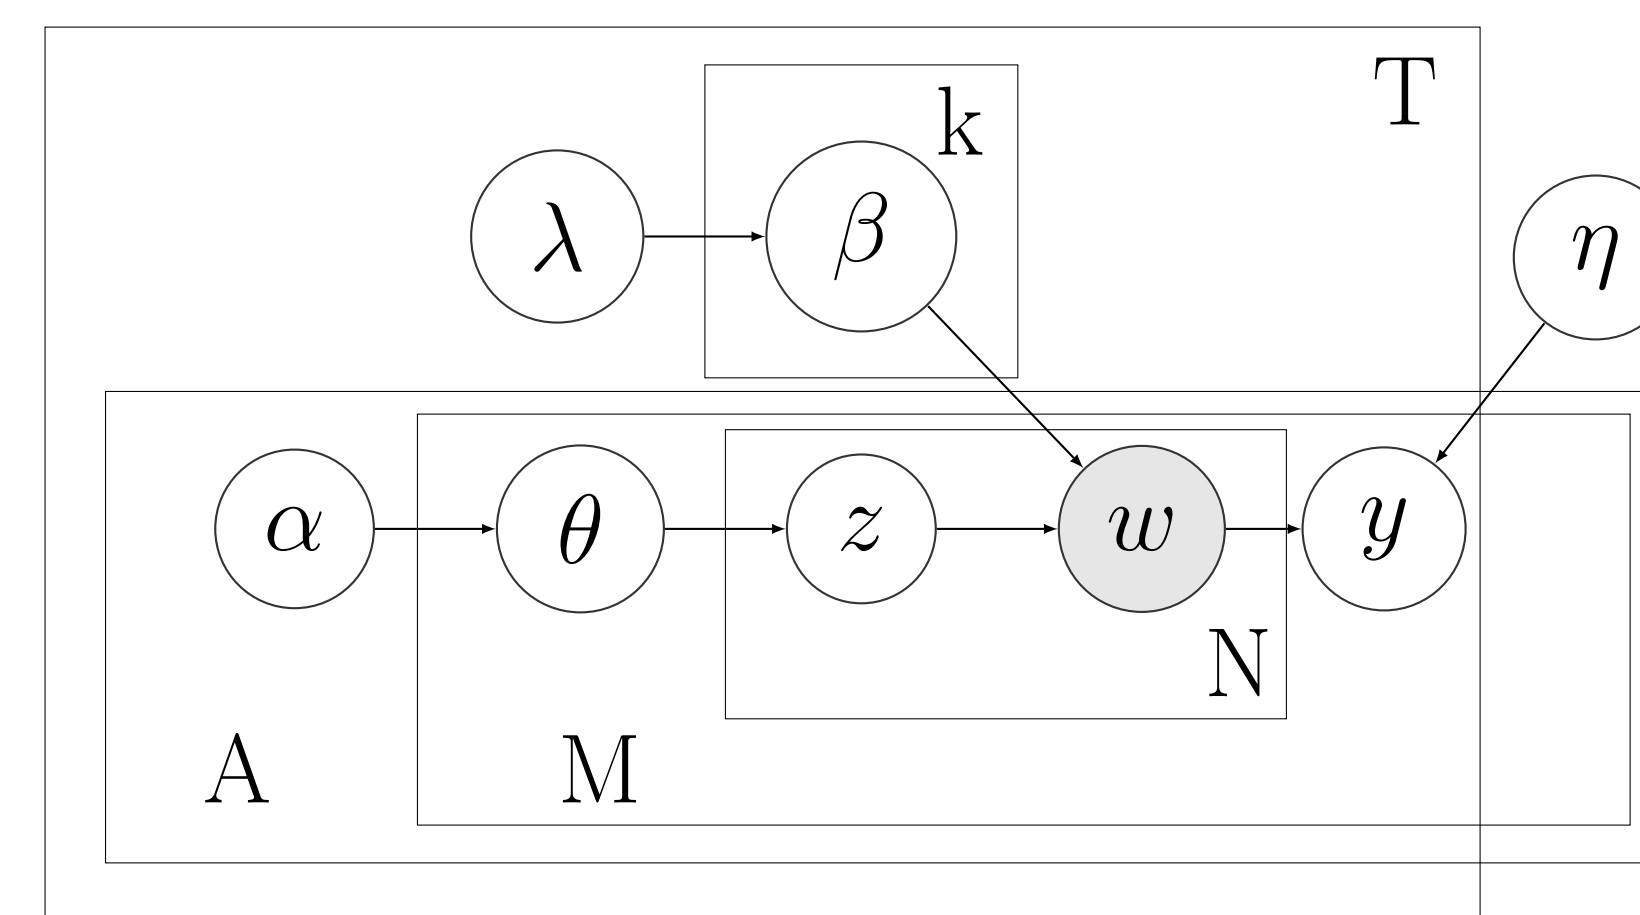
- [1] David M Blei and Jon D McAuliffe. Supervised Topic Models. In *NIPS*, volume 7, pages 121–128, 2007.
- [2] David M Blei, Andrew Y Ng, and Michael I Jordan. Latent Dirichlet Allocation. *Journal of Machine Learning Research*, 3:993–1022, 2003.



Summary

The Multivalence Authorship Detection (MAD) Topic Model extends Latent Dirichlet Allocation [2] to identify authorship in documents with many separate types (“multivalent”) of count features. MAD is “doubly supervised”—it includes a multi-class logistic regression as in [1]—and also fits per-author Dirichlet distributions for each feature type. We test the MAD Topic Model on several real world corpora using a variety of n -gram features, including part-of-speech, syllable stress, and sequences of word lengths.

Model



Graphical Model for the MAD Topic Model

Results

Our results.